

```

1  /* This file contains the main program of MINIX as well as its shutdown code.
2  * The routine main() initializes the system and starts the ball rolling by
3  * setting up the process table, interrupt vectors, and scheduling each task
4  * to run to initialize itself.
5  * The routine shutdown() does the opposite and brings down MINIX.
6  *
7  * The entries into this file are:
8  *   main:             MINIX main program
9  *   prepare_shutdown: prepare to take MINIX down
10 /*
11 #include "kernel.h"
12 #include <signal.h>
13 #include <string.h>
14 #include <unistd.h>
15 #include <a.out.h>
16 #include <minix/callnr.h>
17 #include <minix/com.h>
18 #include <minix/endpoint.h>
19 #include "proc.h"
20
21 /* Prototype declarations for PRIVATE functions. */
22 FORWARD _PROTOTYPE( void announce, (void));
23 FORWARD _PROTOTYPE( void shutdown, (timer_t *));
24
25 /*=====*
26 *                               main                               *
27 *=====*/
28 PUBLIC void main()
29 {
30 /* Start the ball rolling. */
31 struct boot_image *ip;      /* boot image pointer */
32 register struct proc *rp;    /* process pointer */
33 register struct priv *sp;    /* privilege structure pointer */
34 register int i, s;
35 int hdrindex;               /* index to array of a.out headers */
36 phys_clicks text_base;
37 vir_clicks text_clicks, data_clicks, st_clicks;
38 reg_t ktsb;                 /* kernel task stack base */
39 struct exec e_hdr;          /* for a copy of an a.out header */
40
41 /* Clear the process table. Anounce each slot as empty and set up mappings
42 * for proc_addr() and proc_nr() macros. Do the same for the table with
43 * privilege structures for the system processes.
44 */
45 for (rp = BEG_PROC_ADDR, i = -NR_TASKS; rp < END_PROC_ADDR; ++rp, ++i) {
46     rp->p_rts_flags = SLOT_FREE;      /* initialize free slot */
47     rp->p_nr = i;                     /* proc number from ptr */
48     rp->p_endpoint = _ENDPOINT(0, rp->p_nr); /* generation no. 0 */
49     (pproc_addr + NR_TASKS)[i] = rp;  /* proc ptr from number */
50     memset(rp->p_mess_sent, 0, sizeof(rp->p_mess_sent)); /* sent message counter */
51 }
52 for (sp = BEG_PRIV_ADDR, i = 0; sp < END_PRIV_ADDR; ++sp, ++i) {
53     sp->s_proc_nr = NONE;              /* initialize as free */
54     sp->s_id = i;                     /* priv structure index */
55     ppriv_addr[i] = sp;               /* priv ptr from number */
56 }
57
58 /* Set up proc table entries for processes in boot image. The stacks of the
59 * kernel tasks are initialized to an array in data space. The stacks
60 * of the servers have been added to the data segment by the monitor, so
61 * the stack pointer is set to the end of the data segment. All the
62 * processes are in low memory on the 8086. On the 386 only the kernel
63 * is in low memory, the rest is loaded in extended memory.
64 */
65

```

```

66  /* Task stacks. */
67  ktsb = (reg_t) t_stack;
68
69  for (i=0; i < NR_BOOT_PROCS; ++i) {
70      int ci;
71      bitchunk_t fv;
72      ip = &image[i];                /* process' attributes */
73      rp = proc_addr(ip->proc_nr);    /* get process pointer */
74      ip->endpoint = rp->p_endpoint;    /* ipc endpoint */
75      rp->p_max_priority = ip->priority; /* max scheduling priority */
76      rp->p_priority = ip->priority;    /* current priority */
77      rp->p_quantum_size = ip->quantum; /* quantum size in ticks */
78      rp->p_ticks_left = ip->quantum;   /* current credit */
79      strncpy(rp->p_name, ip->proc_name, P_NAME_LEN); /* set process name */
80      (void) get_priv(rp, (ip->flags & SYS_PROC));    /* assign structure */
81      priv(rp)->s_flags = ip->flags;                /* process flags */
82      priv(rp)->s_trap_mask = ip->trap_mask;         /* allowed traps */
83
84      /* Initialize call mask bitmap from unordered set.
85       * A single SYS_ALL_CALLS is a special case - it
86       * means all calls are allowed.
87       */
88      if(ip->nr_k_calls == 1 && ip->k_calls[0] == SYS_ALL_CALLS)
89          fv = ~0;                /* fill call mask */
90      else
91          fv = 0;                /* clear call mask */
92
93      for(ci = 0; ci < CALL_MASK_SIZE; ci++) /* fill or clear call mask */
94          priv(rp)->s_k_call_mask[ci] = fv;
95      if(!fv) /* not all full? enter calls bit by bit */
96          for(ci = 0; ci < ip->nr_k_calls; ci++)
97              SET_BIT(priv(rp)->s_k_call_mask,
98                      ip->k_calls[ci]-KERNEL_CALL);
99
100     priv(rp)->s_ipc_to.chunk[0] = ip->ipc_to; /* restrict targets */
101     if (iskerneln(proc_nr(rp))) { /* part of the kernel? */
102         if (ip->stksize > 0) { /* HARDWARE stack size is 0 */
103             rp->p_priv->s_stack_guard = (reg_t *) ktsb;
104             *rp->p_priv->s_stack_guard = STACK_GUARD;
105         }
106         ktsb += ip->stksize; /* point to high end of stack */
107         rp->p_reg.sp = ktsb; /* this task's initial stack ptr */
108         hdrindex = 0; /* all use the first a.out header */
109     } else {
110         hdrindex = 1 + i-NR_TASKS; /* servers, drivers, INIT */
111     }
112
113     /* The bootstrap loader created an array of the a.out headers at
114      * absolute address 'aout'. Get one element to e_hdr.
115      */
116     phys_copy(aout + hdrindex * A_MINHDR, vir2phys(&e_hdr),
117              (phys_bytes) A_MINHDR);
118     /* Convert addresses to clicks and build process memory map */
119     text_base = e_hdr.a_syms >> CLICK_SHIFT;
120     text_clicks = (e_hdr.a_text + CLICK_SIZE-1) >> CLICK_SHIFT;
121     data_clicks = (e_hdr.a_data+e_hdr.a_bss + CLICK_SIZE-1) >> CLICK_SHIFT;
122     st_clicks = (e_hdr.a_total + CLICK_SIZE-1) >> CLICK_SHIFT;
123     if (!(e_hdr.a_flags & A_SEP))
124     {
125         data_clicks = (e_hdr.a_text+e_hdr.a_data+e_hdr.a_bss +
126                      CLICK_SIZE-1) >> CLICK_SHIFT;
127         text_clicks = 0; /* common I&D */
128     }
129     rp->p_memmap[T].mem_phys = text_base;
130     rp->p_memmap[T].mem_len = text_clicks;

```

```

131     rp->p_memmap[D].mem_phys = text_base + text_clicks;
132     rp->p_memmap[D].mem_len  = data_clicks;
133     rp->p_memmap[S].mem_phys = text_base + text_clicks + st_clicks;
134     rp->p_memmap[S].mem_vir  = st_clicks;
135     rp->p_memmap[S].mem_len  = 0;
136
137     /* Set initial register values. The processor status word for tasks
138      * is different from that of other processes because tasks can
139      * access I/O; this is not allowed to less-privileged processes
140      */
141     rp->p_reg.pc = (reg_t) ip->initial_pc;
142     rp->p_reg.psw = (iskernelp(rp)) ? INIT_TASK_PSW : INIT_PSW;
143
144     /* Initialize the server stack pointer. Take it down one word
145      * to give crtso.s something to use as "argc".
146      */
147     if (isusern(proc_nr(rp))) { /* user-space process? */
148         rp->p_reg.sp = (rp->p_memmap[S].mem_vir +
149                     rp->p_memmap[S].mem_len) << CLICK_SHIFT;
150         rp->p_reg.sp -= sizeof(reg_t);
151     }
152
153     /* Set ready. The HARDWARE task is never ready. */
154     if (rp->p_nr == HARDWARE) RTS_LOCK_SET(rp, NO_PRIORITY);
155     RTS_LOCK_UNSET(rp, SLOT_FREE); /* remove SLOT_FREE and schedule */
156
157     /* Code and data segments must be allocated in protected mode. */
158     alloc_segments(rp);
159 }
160
161 #if SPROFILE
162     sprofiling = 0; /* we're not profiling until instructed to */
163 #endif /* SPROFILE */
164 #if CPROFILE
165     cprof_procs_no = 0; /* init nr of hash table slots used */
166 #endif /* CPROFILE */
167
168     /* MINIX is now ready. All boot image processes are on the ready queue.
169      * Return to the assembly code to start running the current process.
170      */
171     bill_ptr = proc_addr(IDLE); /* it has to point somewhere */
172     announce(); /* print MINIX startup banner */
173     restart();
174 }
175
176 /*=====
177  *                               announce                               *
178  *=====*/
179 PRIVATE void announce(void)
180 {
181     /* Display the MINIX startup banner. */
182     kprintf("\nMINIX %s.%s. "
183 #ifdef _SVN_REVISION
184         "(" _SVN_REVISION ") \n"
185 #endif
186         "Copyright 2006, Vrije Universiteit, Amsterdam, The Netherlands\n",
187         OS_RELEASE, OS_VERSION);
188 }
189
190 /*=====
191  *                               prepare_shutdown                       *
192  *=====*/
193 PUBLIC void prepare_shutdown(how)
194 int how;
195 {

```

```

196  /* This function prepares to shutdown MINIX. */
197  static timer_t shutdown_timer;
198  register struct proc *rp;
199  message m;
200
201  /* Send a signal to all system processes that are still alive to inform
202   * them that the MINIX kernel is shutting down. A proper shutdown sequence
203   * should be implemented by a user-space server. This mechanism is useful
204   * as a backup in case of system panics, so that system processes can still
205   * run their shutdown code, e.g, to synchronize the FS or to let the TTY
206   * switch to the first console.
207   */
208  #if DEAD_CODE
209  kprintf("Sending SIGKSTOP to system processes ...\n");
210  for (rp=BEG_PROC_ADDR; rp<END_PROC_ADDR; rp++) {
211      if (!isemptyp(rp) && (priv(rp)->s_flags & SYS_PROC) && !iskernelp(rp))
212          send_sig(proc_nr(rp), SIGKSTOP);
213  }
214  #endif
215
216  /* Continue after 1 second, to give processes a chance to get scheduled to
217   * do shutdown work. Set a watchdog timer to call shutdown(). The timer
218   * argument passes the shutdown status.
219   */
220  kprintf("MINIX will now be shut down ...\n");
221  tmr_arg(&shutdown_timer)->ta_int = how;
222  set_timer(&shutdown_timer, get_uptime() + HZ, shutdown);
223  }
224
225  /*=====*
226   *                               shutdown                               *
227   *=====*/
228  PRIVATE void shutdown(tp)
229  timer_t *tp;
230  {
231  /* This function is called from prepare_shutdown or stop_sequence to bring
232   * down MINIX. How to shutdown is in the argument: RBT_HALT (return to the
233   * monitor), RBT_MONITOR (execute given code), RBT_RESET (hard reset).
234   */
235  intr_init(INTS_ORIG);
236  clock_stop();
237  arch_shutdown(tmr_arg(tp)->ta_int);
238  }
239
240

```