

```

1  /* This file contains essentially all of the process and message handling.
2  * Together with "mpx.s" it forms the lowest layer of the MINIX kernel.
3  * There is one entry point from the outside:
4  *
5  *   sys_call:      a system call, i.e., the kernel is trapped with an INT
6  *
7  * As well as several entry points used from the interrupt and task level:
8  *
9  *   lock_notify:    notify a process of a system event
10 *   lock_send:      send a message to a process
11 *   lock_enqueue:   put a process on one of the scheduling queues
12 *   lock_dequeue:   remove a process from the scheduling queues
13 *
14 * Changes:
15 *   Aug 19, 2005      rewrote scheduling code (Jorrit N. Herder)
16 *   Jul 25, 2005      rewrote system call handling (Jorrit N. Herder)
17 *   May 26, 2005      rewrote message passing functions (Jorrit N. Herder)
18 *   May 24, 2005      new notification system call (Jorrit N. Herder)
19 *   Oct 28, 2004      nonblocking send and receive calls (Jorrit N. Herder)
20 *
21 * The code here is critical to make everything work and is important for the
22 * overall performance of the system. A large fraction of the code deals with
23 * list manipulation. To make this both easy to understand and fast to execute
24 * pointer pointers are used throughout the code. Pointer pointers prevent
25 * exceptions for the head or tail of a linked list.
26 *
27 *   node_t *queue, *new_node;    // assume these as global variables
28 *   node_t **xpp = &queue;      // get pointer pointer to head of queue
29 *   while (*xpp != NULL)        // find last pointer of the linked list
30 *       xpp = &(*xpp)->next;    // get pointer to next pointer
31 *   *xpp = new_node;            // now replace the end (the NULL pointer)
32 *   new_node->next = NULL;      // and mark the new end of the list
33 *
34 * For example, when adding a new node to the end of the list, one normally
35 * makes an exception for an empty list and looks up the end of the list for
36 * nonempty lists. As shown above, this is not required with pointer pointers.
37 */
38
39 #include <minix/com.h>
40 #include <minix/callnr.h>
41 #include <minix/endpoint.h>
42 #include "debug.h"
43 #include "kernel.h"
44 #include "proc.h"
45 #include <signal.h>
46 #include <minix/portio.h>
47
48 /* Scheduling and message passing functions. The functions are available to
49 * other parts of the kernel through lock_...(). The lock temporarily disables
50 * interrupts to prevent race conditions.
51 */
52 FORWARD _PROTOTYPE( int mini_send, (struct proc *caller_ptr, int dst_e,
53                                   message *m_ptr, unsigned flags));
54 FORWARD _PROTOTYPE( int mini_receive, (struct proc *caller_ptr, int src,
55                                       message *m_ptr, unsigned flags));
56 FORWARD _PROTOTYPE( int mini_notify, (struct proc *caller_ptr, int dst));
57 FORWARD _PROTOTYPE( int deadlock, (int function,
58                                   register struct proc *caller, int src_dst));
59 FORWARD _PROTOTYPE( void enqueue, (struct proc *rp));
60 FORWARD _PROTOTYPE( void dequeue, (struct proc *rp));
61 FORWARD _PROTOTYPE( void sched, (struct proc *rp, int *queue, int *front));
62 FORWARD _PROTOTYPE( void pick_proc, (void));
63
64 #define BuildMess(m_ptr, src, dst_ptr) \
65     (m_ptr)->m_source = proc_addr(src)->p_endpoint; \

```

```

66         (m_ptr)->m_type = NOTIFY_FROM(src); \
67         (m_ptr)->NOTIFY_TIMESTAMP = get_uptime(); \
68         switch (src) { \
69             case HARDWARE: \
70                 (m_ptr)->NOTIFY_ARG = priv(dst_ptr)->s_int_pending; \
71                 priv(dst_ptr)->s_int_pending = 0; \
72                 break; \
73             case SYSTEM: \
74                 (m_ptr)->NOTIFY_ARG = priv(dst_ptr)->s_sig_pending; \
75                 priv(dst_ptr)->s_sig_pending = 0; \
76                 break; \
77         }
78
79 #define CopyMess(s,sp,sm,dp,dm) \
80     cp_mess(proc_addr(s)->p_endpoint, \
81             (sp)->p_memmap[D].mem_phys, \
82             (vir_bytes)sm, (dp)->p_memmap[D].mem_phys, (vir_bytes)dm)
83
84 /*=====
85  *                               sys_call                               *
86  *=====*/
87 PUBLIC int sys_call(call_nr, src_dst_e, m_ptr, bit_map)
88 int call_nr; /* system call number and flags */
89 int src_dst_e; /* src to receive from or dst to send to */
90 message *m_ptr; /* pointer to message in the caller's space */
91 long bit_map; /* notification event set or flags */
92 {
93     /* System calls are done by trapping to the kernel with an INT instruction.
94      * The trap is caught and sys_call() is called to send or receive a message
95      * (or both). The caller is always given by 'proc_ptr'.
96      */
97     register struct proc *caller_ptr = proc_ptr; /* get pointer to caller */
98     int function = call_nr & SYSCALL_FUNC; /* get system call function */
99     unsigned flags = call_nr & SYSCALL_FLAGS; /* get flags */
100     int mask_entry; /* bit to check in send mask */
101     int group_size; /* used for deadlock check */
102     int result; /* the system call's result */
103     int src_dst;
104     vir_clicks vlo, vhi; /* virtual clicks containing message to send */
105
106     #if 1
107         if (RTS_ISSET(caller_ptr, SLOT_FREE))
108         {
109             kprintf("called by the dead?!?\n");
110             return EINVAL;
111         }
112     #endif
113
114     /* Require a valid source and/ or destination process, unless echoing. */
115     if (src_dst_e != ANY && function != ECHO) {
116         if (!isokendpt(src_dst_e, &src_dst)) {
117             #if DEBUG_ENABLE_IPC_WARNINGS
118                 kprintf("sys_call: trap %d by %d with bad endpoint %d\n",
119                     function, proc_nr(caller_ptr), src_dst_e);
120             #endif
121             return EDEADSRCDST;
122         }
123     } else src_dst = src_dst_e;
124
125     /* Check if the process has privileges for the requested call. Calls to the
126      * kernel may only be SENDREC, because tasks always reply and may not block
127      * if the caller doesn't do receive().
128      */
129     if (! (priv(caller_ptr)->s_trap_mask & (1 << function)) ||
130         (iskerneln(src_dst) && function != SENDREC

```

```

131         && function != RECEIVE)) {
132     #if DEBUG_ENABLE_IPC_WARNINGS
133         kprintf("sys_call: trap %d not allowed, caller %d, src_dst %d\n",
134             function, proc_nr(caller_ptr), src_dst);
135     #endif
136     return(ETRAPDENIED);          /* trap denied by mask or kernel */
137 }
138
139 /* If the call involves a message buffer, i.e., for SEND, RECEIVE, SENDREC,
140  * or ECHO, check the message pointer. This check allows a message to be
141  * anywhere in data or stack or gap. It will have to be made more elaborate
142  * for machines which don't have the gap mapped.
143  */
144 if (function & CHECK_PTR) {
145     vlo = (vir_bytes) m_ptr >> CLICK_SHIFT;
146     vhi = ((vir_bytes) m_ptr + MESS_SIZE - 1) >> CLICK_SHIFT;
147     if (vlo < caller_ptr->p_memmap[D].mem_vir || vlo > vhi ||
148         vhi >= caller_ptr->p_memmap[S].mem_vir +
149             caller_ptr->p_memmap[S].mem_len) {
150 #if DEBUG_ENABLE_IPC_WARNINGS
151     kprintf("sys_call: invalid message pointer, trap %d, caller %d\n",
152         function, proc_nr(caller_ptr));
153 #endif
154     return(EFAULT);              /* invalid message pointer */
155 }
156 }
157
158 /* If the call is to send to a process, i.e., for SEND, SENDREC or NOTIFY,
159  * verify that the caller is allowed to send to the given destination.
160  */
161 if (function & CHECK_DST) {
162     if (! get_sys_bit(priv(caller_ptr)->s_ipc_to, nr_to_id(src_dst))) {
163 #if DEBUG_ENABLE_IPC_WARNINGS
164     kprintf("sys_call: ipc mask denied trap %d from %d to %d\n",
165         function, proc_nr(caller_ptr), src_dst);
166 #endif
167     return(ECALLDENIED);          /* call denied by ipc mask */
168 }
169 }
170
171 /* Check for a possible deadlock for blocking SEND(REC) and RECEIVE. */
172 if (function & CHECK_DEADLOCK) {
173     if (group_size = deadlock(function, caller_ptr, src_dst)) {
174 #if DEBUG_ENABLE_IPC_WARNINGS
175     kprintf("sys_call: trap %d from %d to %d deadlocked, group size %d\n",
176         function, proc_nr(caller_ptr), src_dst, group_size);
177 #endif
178     return(ELCKED);
179 }
180 }
181
182 /* Now check if the call is known and try to perform the request. The only
183  * system calls that exist in MINIX are sending and receiving messages.
184  * - SENDREC: combines SEND and RECEIVE in a single system call
185  * - SEND: sender blocks until its message has been delivered
186  * - RECEIVE: receiver blocks until an acceptable message has arrived
187  * - NOTIFY: nonblocking call; deliver notification or mark pending
188  * - ECHO: nonblocking call; directly echo back the message
189  */
190 switch(function) {
191 case SENDREC:
192     /* A flag is set so that notifications cannot interrupt SENDREC. */
193     caller_ptr->p_misc_flags |= REPLY_PENDING;
194     /* fall through */
195 case SEND:

```

```

196     result = mini_send(caller_ptr, src_dst_e, m_ptr, flags);
197     if (function == SEND || result != OK) {
198         break; /* done, or SEND failed */
199     } /* fall through for SENDREC */
200 case RECEIVE:
201     if (function == RECEIVE)
202         caller_ptr->p_misc_flags &= ~REPLY_PENDING;
203     result = mini_receive(caller_ptr, src_dst_e, m_ptr, flags);
204     break;
205 case NOTIFY:
206     result = mini_notify(caller_ptr, src_dst);
207     break;
208 case ECHO:
209     CopyMess(caller_ptr->p_nr, caller_ptr, m_ptr, caller_ptr, m_ptr);
210     result = OK;
211     break;
212 default:
213     result = EBADCALL; /* illegal system call */
214 }
215
216 /* Now, return the result of the system call to the caller. */
217 return(result);
218 }
219
220 /*=====
221 *                                deadlock                                *
222 *=====*/
223 PRIVATE int deadlock(function, cp, src_dst)
224 int function; /* trap number */
225 register struct proc *cp; /* pointer to caller */
226 int src_dst; /* src or dst process */
227 {
228 /* Check for deadlock. This can happen if 'caller_ptr' and 'src_dst' have
229 * a cyclic dependency of blocking send and receive calls. The only cyclic
230 * dependency that is not fatal is if the caller and target directly SEND(REC)
231 * and RECEIVE to each other. If a deadlock is found, the group size is
232 * returned. Otherwise zero is returned.
233 */
234 register struct proc *xp; /* process pointer */
235 int group_size = 1; /* start with only caller */
236 int trap_flags;
237
238 while (src_dst != ANY) { /* check while process nr */
239     int src_dst_e;
240     xp = proc_addr(src_dst); /* follow chain of processes */
241     group_size++; /* extra process in group */
242
243     /* Check whether the last process in the chain has a dependency. If it
244     * has not, the cycle cannot be closed and we are done.
245     */
246     if (RTS_ISSET(xp, RECEIVING)) { /* xp has dependency */
247         if (xp->p_getfrom_e == ANY) src_dst = ANY;
248         else okendpt(xp->p_getfrom_e, &src_dst);
249     } else if (RTS_ISSET(xp, SENDING)) { /* xp has dependency */
250         okendpt(xp->p_sendto_e, &src_dst);
251     } else {
252         return(0); /* not a deadlock */
253     }
254
255     /* Now check if there is a cyclic dependency. For group sizes of two,
256     * a combination of SEND(REC) and RECEIVE is not fatal. Larger groups
257     * or other combinations indicate a deadlock.
258     */
259     if (src_dst == proc_nr(cp)) { /* possible deadlock */
260         if (group_size == 2) { /* caller and src_dst */

```

```

261         /* The function number is magically converted to flags. */
262         if ((xp->p_rts_flags ^ (function << 2)) & SENDING) {
263             return(0); /* not a deadlock */
264         }
265     }
266     return(group_size); /* deadlock found */
267 }
268 }
269 return(0); /* not a deadlock */
270 }
271
272 /*=====
273      * mini_send
274      *=====*/
275 PRIVATE int mini_send(caller_ptr, dst_e, m_ptr, flags)
276 register struct proc *caller_ptr; /* who is trying to send a message? */
277 int dst_e; /* to whom is message being sent? */
278 message *m_ptr; /* pointer to message buffer */
279 unsigned flags; /* system call flags */
280 {
281     /* Send a message from 'caller_ptr' to 'dst'. If 'dst' is blocked waiting
282      * for this message, copy the message to it and unblock 'dst'. If 'dst' is
283      * not waiting at all, or is waiting for another source, queue 'caller_ptr'.
284      */
285     register struct proc *dst_ptr;
286     register struct proc **xpp;
287     int dst_p;
288
289     dst_p = _ENDPOINT_P(dst_e);
290     dst_ptr = proc_addr(dst_p);
291
292     if (RTS_ISSET(dst_ptr, NO_ENDPOINT)) return EDSTDIED;
293
294     /* Check if 'dst' is blocked waiting for this message. The destination's
295      * SENDING flag may be set when its SENDREC call blocked while sending.
296      */
297     if ( (RTS_ISSET(dst_ptr, RECEIVING) && !RTS_ISSET(dst_ptr, SENDING)) &&
298          (dst_ptr->p_getfrom_e == ANY
299           || dst_ptr->p_getfrom_e == caller_ptr->p_endpoint)) {
300         /* Destination is indeed waiting for this message. */
301         CopyMess(caller_ptr->p_nr, caller_ptr, m_ptr, dst_ptr,
302                  dst_ptr->p_messbuf);
303         RTS_UNSET(dst_ptr, RECEIVING);
304     } else if ( ! (flags & NON_BLOCKING)) {
305         /* Destination is not waiting. Block and dequeue caller. */
306         caller_ptr->p_messbuf = m_ptr;
307         RTS_SET(caller_ptr, SENDING);
308         caller_ptr->p_sendto_e = dst_e;
309
310         /* Process is now blocked. Put in on the destination's queue. */
311         xpp = &dst_ptr->p_caller_q; /* find end of list */
312         while (*xpp != NIL_PROC) xpp = &(*xpp)->p_q_link;
313         *xpp = caller_ptr; /* add caller to end */
314         caller_ptr->p_q_link = NIL_PROC; /* mark new end of list */
315     } else {
316         return(ENOTREADY);
317     }
318
319     /* Increment the counter keeping track of where messages are sent. */
320     ++(caller_ptr->p_mess_sent[dst_ptr->p_nr + NR_TASKS]);
321
322     return(OK);
323 }
324
325 /*=====

```

```

326      *                                mini_receive                                *
327      *=====*/
328  PRIVATE int mini_receive(caller_ptr, src_e, m_ptr, flags)
329  register struct proc *caller_ptr;      /* process trying to get message */
330  int src_e;                             /* which message source is wanted */
331  message *m_ptr;                        /* pointer to message buffer */
332  unsigned flags;                        /* system call flags */
333  {
334  /* A process or task wants to get a message.  If a message is already queued,
335  * acquire it and deblock the sender.  If no message from the desired source
336  * is available block the caller, unless the flags don't allow blocking.
337  */
338  register struct proc **xpp;
339  register struct notification **ntf_q_pp;
340  message m;
341  int bit_nr;
342  sys_map_t *map;
343  bitchunk_t *chunk;
344  int i, src_id, src_proc_nr, src_p;
345
346  if(src_e == ANY) src_p = ANY;
347  else
348  {
349      okendpt(src_e, &src_p);
350      if (RTS_ISSET(proc_addr(src_p), NO_ENDPOINT)) return ESRCDIED;
351  }
352
353
354  /* Check to see if a message from desired source is already available.
355  * The caller's SENDING flag may be set if SENDREC couldn't send. If it is
356  * set, the process should be blocked.
357  */
358  if (!RTS_ISSET(caller_ptr, SENDING)) {
359
360      /* Check if there are pending notifications, except for SENDREC. */
361      if (! (caller_ptr->p_misc_flags & REPLY_PENDING)) {
362
363          map = &priv(caller_ptr)->s_notify_pending;
364          for (chunk=&map->chunk[0]; chunk<&map->chunk[NR_SYS_CHUNKS]; chunk++) {
365
366              /* Find a pending notification from the requested source. */
367              if (! *chunk) continue; /* no bits in chunk */
368              for (i=0; ! (*chunk & (1<<i)); ++i) {} /* look up the bit */
369              src_id = (chunk - &map->chunk[0]) * BITCHUNK_BITS + i;
370              if (src_id >= NR_SYS_PROCS) break; /* out of range */
371              src_proc_nr = id_to_nr(src_id); /* get source proc */
372
373              #if DEBUG_ENABLE_IPC_WARNINGS
374              if(src_proc_nr == NONE) {
375                  kprintf("mini_receive: sending notify from NONE\n");
376              }
377              #endif
378
379              if (src_e!=ANY && src_p != src_proc_nr) continue; /* source not ok */
380              *chunk &= ~(1 << i); /* no longer pending */
381
382              /* Found a suitable source, deliver the notification message. */
383              BuildMess(&m, src_proc_nr, caller_ptr); /* assemble message */
384              CopyMess(src_proc_nr, proc_addr(HARDWARE), &m, caller_ptr, m_ptr);
385              return(OK); /* report success */
386          }
387      }
388
389      /* Check caller queue. Use pointer pointers to keep code simple. */
390      xpp = &caller_ptr->p_caller_q;
391      while (*xpp != NIL_PROC) {
392          if (src_e == ANY || src_p == proc_nr(*xpp)) {

```

```

391     #if 1
392         if (RTS_ISSET(*xpp, SLOT_FREE))
393         {
394             kprintf("listening to the dead?!?\n");
395             return EINVAL;
396         }
397     #endif
398
399     /* Found acceptable message. Copy it and update status. */
400     CopyMess((*xpp)->p_nr, *xpp, (*xpp)->p_messbuf, caller_ptr, m_ptr);
401     RTS_UNSET(*xpp, SENDING);
402     *xpp = (*xpp)->p_q_link;          /* remove from queue */
403     return(OK);                      /* report success */
404 }
405 xpp = &(*xpp)->p_q_link;          /* proceed to next */
406 }
407 }
408
409 /* No suitable message is available or the caller couldn't send in SENDREC.
410  * Block the process trying to receive, unless the flags tell otherwise.
411  */
412 if ( ! (flags & NON_BLOCKING)) {
413     caller_ptr->p_getfrom_e = src_e;
414     caller_ptr->p_messbuf = m_ptr;
415     RTS_SET(caller_ptr, RECEIVING);
416     return(OK);
417 } else {
418     return(ENOTREADY);
419 }
420 }
421
422 /*=====
423  *                               mini_notify                               *
424  *=====*/
425 PRIVATE int mini_notify(caller_ptr, dst)
426 register struct proc *caller_ptr;    /* sender of the notification */
427 int dst;                            /* which process to notify */
428 {
429     register struct proc *dst_ptr = proc_addr(dst);
430     int src_id;                      /* source id for late delivery */
431     message m;                      /* the notification message */
432
433     /* Check to see if target is blocked waiting for this message. A process
434      * can be both sending and receiving during a SENDREC system call.
435      */
436     if ( (RTS_ISSET(dst_ptr, RECEIVING) && !RTS_ISSET(dst_ptr, SENDING)) &&
437          ! (dst_ptr->p_misc_flags & REPLY_PENDING) &&
438          (dst_ptr->p_getfrom_e == ANY ||
439           dst_ptr->p_getfrom_e == caller_ptr->p_endpoint)) {
440
441         /* Destination is indeed waiting for a message. Assemble a notification
442          * message and deliver it. Copy from pseudo-source HARDWARE, since the
443          * message is in the kernel's address space.
444          */
445         BuildMess(&m, proc_nr(caller_ptr), dst_ptr);
446         CopyMess(proc_nr(caller_ptr), proc_addr(HARDWARE), &m,
447                 dst_ptr, dst_ptr->p_messbuf);
448         RTS_UNSET(dst_ptr, RECEIVING);
449         return(OK);
450     }
451
452     /* Destination is not ready to receive the notification. Add it to the
453      * bit map with pending notifications. Note the indirectness: the system id
454      * instead of the process number is used in the pending bit map.
455      */

```



```

521     rp->p_nextready = rdy_head[q];          /* chain head of queue */
522     rdy_head[q] = rp;                      /* set new queue head */
523 }
524 else {                                     /* add to tail of queue */
525     rdy_tail[q]->p_nextready = rp;          /* chain tail of queue */
526     rdy_tail[q] = rp;                     /* set new queue tail */
527     rp->p_nextready = NIL_PROC;             /* mark new end */
528 }
529
530 /* Now select the next process to run, if there isn't a current
531  * process yet or current process isn't ready any more, or
532  * it's PREEMPTIBLE.
533  */
534 if(!proc_ptr || proc_ptr->p_rts_flags ||
535     (priv(proc_ptr)->s_flags & PREEMPTIBLE)) {
536     pick_proc();
537 }
538
539 #if DEBUG_SCHED_CHECK
540     rp->p_ready = 1;
541     check_runqueues("enqueue2");
542 #endif
543 }
544
545 /*=====
546  *                               dequeue                               *
547  *=====*/
548 PRIVATE void dequeue(rp)
549 register struct proc *rp;          /* this process is no longer runnable */
550 {
551     /* A process must be removed from the scheduling queues, for example, because
552      * it has blocked. If the currently active process is removed, a new process
553      * is picked to run by calling pick_proc().
554      */
555     register int q = rp->p_priority;      /* queue to use */
556     register struct proc **xpp;          /* iterate over queue */
557     register struct proc *prev_xp;
558
559     /* Side-effect for kernel: check if the task's stack still is ok? */
560     if (iskernelp(rp)) {
561         if (*priv(rp)->s_stack_guard != STACK_GUARD)
562             panic("stack overrun by task", proc_nr(rp));
563     }
564
565     #if DEBUG_SCHED_CHECK
566         check_runqueues("dequeue1");
567         if (!rp->p_ready) kprintf("dequeue() already unready process\n");
568     #endif
569
570     /* Now make sure that the process is not in its ready queue. Remove the
571      * process if it is found. A process can be made unready even if it is not
572      * running by being sent a signal that kills it.
573      */
574     prev_xp = NIL_PROC;
575     for (xpp = &rdy_head[q]; *xpp != NIL_PROC; xpp = &(*xpp)->p_nextready) {
576
577         if (*xpp == rp) {                 /* found process to remove */
578             *xpp = (*xpp)->p_nextready;   /* replace with next chain */
579             if (rp == rdy_tail[q])        /* queue tail removed */
580                 rdy_tail[q] = prev_xp;    /* set new tail */
581             if (rp == proc_ptr || rp == next_ptr) /* active process removed */
582                 pick_proc();              /* pick new process to run */
583             break;
584         }
585         prev_xp = *xpp;                  /* save previous in chain */

```

```

586     }
587
588     #if DEBUG_SCHED_CHECK
589         rp->p_ready = 0;
590         check_runqueues("dequeue2");
591     #endif
592 }
593
594 /*=====
595      *                                sched                                *
596      *=====*/
597 PRIVATE void sched(rp, queue, front)
598     register struct proc *rp;          /* process to be scheduled */
599     int *queue;                        /* return: queue to use */
600     int *front;                        /* return: front or back */
601 {
602     /* This function determines the scheduling policy.  It is called whenever a
603      * process must be added to one of the scheduling queues to decide where to
604      * insert it.  As a side-effect the process' priority may be updated.
605      */
606     int time_left = (rp->p_ticks_left > 0);    /* quantum fully consumed */
607
608     /* Check whether the process has time left.  Otherwise give a new quantum
609      * and lower the process' priority, unless the process already is in the
610      * lowest queue.
611      */
612     if (! time_left) {                    /* quantum consumed ? */
613         rp->p_ticks_left = rp->p_quantum_size; /* give new quantum */
614         if (rp->p_priority < (IDLE_Q-1)) {
615             rp->p_priority += 1;          /* lower priority */
616         }
617     }
618
619     /* If there is time left, the process is added to the front of its queue,
620      * so that it can immediately run.  The queue to use simply is always the
621      * process' current priority.
622      */
623     *queue = rp->p_priority;
624     *front = time_left;
625 }
626
627 /*=====
628      *                                pick_proc                                *
629      *=====*/
630 PRIVATE void pick_proc()
631 {
632     /* Decide who to run now.  A new process is selected by setting 'next_ptr'.
633      * When a billable process is selected, record it in 'bill_ptr', so that the
634      * clock task can tell who to bill for system time.
635      */
636     register struct proc *rp;          /* process to run */
637     int q;                             /* iterate over queues */
638
639     /* Check each of the scheduling queues for ready processes.  The number of
640      * queues is defined in proc.h, and priorities are set in the task table.
641      * The lowest queue contains IDLE, which is always ready.
642      */
643     for (q=0; q < NR_SCHED_QUEUES; q++) {
644         if ( (rp = rdy_head[q]) != NIL_PROC) {
645             next_ptr = rp;              /* run process 'rp' next */
646             if (priv(rp)->s_flags & BILLABLE)
647                 bill_ptr = rp;          /* bill for system time */
648             return;
649         }
650     }

```

```

651     panic("no ready process", NO_NUM);
652 }
653
654 /*=====*
655 *                               balance_queues                               *
656 *=====*/
657 #define Q_BALANCE_TICKS 100
658 PUBLIC void balance_queues(tp)
659 timer_t *tp;                                /* watchdog timer pointer */
660 {
661     /* Check entire process table and give all process a higher priority. This
662      * effectively means giving a new quantum. If a process already is at its
663      * maximum priority, its quantum will be renewed.
664      */
665     static timer_t queue_timer;              /* timer structure to use */
666     register struct proc* rp;                /* process table pointer */
667     clock_t next_period;                     /* time of next period */
668     int ticks_added = 0;                    /* total time added */
669
670     for (rp=BEG_PROC_ADDR; rp<END_PROC_ADDR; rp++) {
671         if (! isemptyp(rp)) {                /* check slot use */
672             lock(5, "balance_queues");
673             if (rp->p_priority > rp->p_max_priority) { /* update priority? */
674                 if (rp->p_rts_flags == 0) dequeue(rp); /* take off queue */
675                 ticks_added += rp->p_quantum_size; /* do accounting */
676                 rp->p_priority -= 1; /* raise priority */
677                 if (rp->p_rts_flags == 0) enqueue(rp); /* put on queue */
678             }
679             else {
680                 ticks_added += rp->p_quantum_size - rp->p_ticks_left;
681                 rp->p_ticks_left = rp->p_quantum_size; /* give new quantum */
682             }
683             unlock(5);
684         }
685     }
686     #if DEBUG
687     kprintf("ticks_added: %d\n", ticks_added);
688     #endif
689
690     /* Now schedule a new watchdog timer to balance the queues again. The
691      * period depends on the total amount of quantum ticks added.
692      */
693     next_period = MAX(Q_BALANCE_TICKS, ticks_added); /* calculate next */
694     set_timer(&queue_timer, get_uptime() + next_period, balance_queues);
695 }
696
697 /*=====*
698 *                               lock_send                               *
699 *=====*/
700 PUBLIC int lock_send(dst_e, m_ptr)
701 int dst_e;                                /* to whom is message being sent? */
702 message *m_ptr;                          /* pointer to message buffer */
703 {
704     /* Safe gateway to mini_send() for tasks. */
705     int result;
706     lock(2, "send");
707     result = mini_send(proc_ptr, dst_e, m_ptr, NON_BLOCKING);
708     unlock(2);
709     return(result);
710 }
711
712 /*=====*
713 *                               lock_enqueue                               *
714 *=====*/
715 PUBLIC void lock_enqueue(rp)

```

```

716 struct proc *rp;                /* this process is now runnable */
717 {
718 /* Safe gateway to enqueue() for tasks. */
719     lock(3, "enqueue");
720     enqueue(rp);
721     unlock(3);
722 }
723
724 /*=====*
725 *                      lock_dequeue                      *
726 *=====*/
727 PUBLIC void lock_dequeue(rp)
728 struct proc *rp;                /* this process is no longer runnable */
729 {
730 /* Safe gateway to dequeue() for tasks. */
731     if (k_reenter >= 0) {
732         /* We're in an exception or interrupt, so don't lock (and ...
733          * don't unlock).
734          */
735         dequeue(rp);
736     } else {
737         lock(4, "dequeue");
738         dequeue(rp);
739         unlock(4);
740     }
741 }
742
743 /*=====*
744 *                      isokendpt_f                      *
745 *=====*/
746 #if DEBUG_ENABLE_IPC_WARNINGS
747 PUBLIC int isokendpt_f(file, line, e, p, fatalflag)
748 char *file;
749 int line;
750 #else
751 PUBLIC int isokendpt_f(e, p, fatalflag)
752 #endif
753 endpoint_t e;
754 int *p, fatalflag;
755 {
756     int ok = 0;
757     /* Convert an endpoint number into a process number.
758      * Return nonzero if the process is alive with the corresponding
759      * generation number, zero otherwise.
760      *
761      * This function is called with file and line number by the
762      * isokendpt_d macro if DEBUG_ENABLE_IPC_WARNINGS is defined,
763      * otherwise without. This allows us to print the where the
764      * conversion was attempted, making the errors verbose without
765      * adding code for that at every call.
766      *
767      * If fatalflag is nonzero, we must panic if the conversion doesn't
768      * succeed.
769      */
770     *p = _ENDPOINT_P(e);
771     if (!isokprocn(*p)) {
772 #if DEBUG_ENABLE_IPC_WARNINGS
773         kprintf("kernel:%s:%d: bad endpoint %d: proc %d out of range\n",
774             file, line, e, *p);
775 #endif
776     } else if (isempty(*p)) {
777 #if DEBUG_ENABLE_IPC_WARNINGS
778         kprintf("kernel:%s:%d: bad endpoint %d: proc %d empty\n", file, line, e, *p);
779 #endif
780     } else if (proc_addr(*p)->p_endpoint != e) {

```

```
781  #if DEBUG_ENABLE_IPC_WARNINGS
782      kprintf("kernel:%s:%d: bad endpoint %d: proc %d has ept %d (generation
783              %d vs. %d)\n", file, line,
784              e, *p, proc_addr(*p)->p_endpoint,
785              _ENDPOINT_G(e), _ENDPOINT_G(proc_addr(*p)->p_endpoint));
786  #endif
787      } else ok = 1;
788      if(!ok && fatalflag) {
789          panic("invalid endpoint ", e);
790      }
791      return ok;
792  }
793
```