

```

1  #ifndef PROC_H
2  #define PROC_H
3
4  /* Here is the declaration of the process table.  It contains all process
5   * data, including registers, flags, scheduling priority, memory map,
6   * accounting, message passing (IPC) information, and so on.
7   *
8   * Many assembly code routines reference fields in it.  The offsets to these
9   * fields are defined in the assembler include file sconst.h.  When changing
10  * struct proc, be sure to change sconst.h to match.
11  */
12  #include <minix/com.h>
13  #include "const.h"
14  #include "priv.h"
15
16  struct proc {
17      struct stackframe_s p_reg;      /* process' registers saved in stack frame */
18      struct segframe p_seg;          /* segment descriptors */
19      proc_nr_t p_nr;                 /* number of this process (for fast access) */
20      struct priv *p_priv;             /* system privileges structure */
21      short p_rts_flags;              /* process is runnable only if zero */
22      short p_misc_flags;             /* flags that do not suspend the process */
23
24      char p_priority;                /* current scheduling priority */
25      char p_max_priority;            /* maximum scheduling priority */
26      char p_ticks_left;              /* number of scheduling ticks left */
27      char p_quantum_size;            /* quantum size in ticks */
28
29      struct mem_map p_memmap[NR_LOCAL_SEGS]; /* memory map (T, D, S) */
30
31      clock_t p_user_time;            /* user time in ticks */
32      clock_t p_sys_time;            /* sys time in ticks */
33
34      struct proc *p_nextready;       /* pointer to next ready process */
35      struct proc *p_caller_q;        /* head of list of procs wishing to send */
36      struct proc *p_q_link;          /* link to next proc wishing to send */
37      message *p_messbuf;             /* pointer to passed message buffer */
38      int p_getfrom_e;                /* from whom does process want to receive? */
39      int p_sendto_e;                /* to whom does process want to send? */
40
41      sigset_t p_pending;             /* bit map for pending kernel signals */
42
43      char p_name[P_NAME_LEN];        /* name of the process, including \0 */
44
45      endpoint_t p_endpoint;          /* endpoint number, generation-aware */
46
47      unsigned long p_mess_sent[NR_TASKS + NR_PROCS]; /* number of messages sent to other
48      processes */
49
50  #if DEBUG_SCHED_CHECK
51      int p_ready, p_found;
52  #endif
53  };
54
55  /* Bits for the runtime flags.  A process is runnable iff p_rts_flags == 0. */
56  #define SLOT_FREE      0x01  /* process slot is free */
57  #define NO_PRIORITY    0x02  /* process has been stopped */
58  #define SENDING        0x04  /* process blocked trying to send */
59  #define RECEIVING      0x08  /* process blocked trying to receive */
60  #define SIGNALLED      0x10  /* set when new kernel signal arrives */
61  #define SIG_PENDING    0x20  /* unready while signal being processed */
62  #define P_STOP         0x40  /* set when process is being traced */
63  #define NO_PRIV        0x80  /* keep forked system process from running */
64  #define NO_ENDPOINT    0x100 /* process cannot send or receive messages */

```

```

65  /* These runtime flags can be tested and manipulated by these macros. */
66
67  #define RTS_ISSET(rp, f) (((rp)->p_rts_flags & (f)) == (f))
68
69
70  /* Set flag and dequeue if the process was runnable. */
71  #define RTS_SET(rp, f) \
72      do { \
73          if(!(rp)->p_rts_flags) { dequeue(rp); } \
74          (rp)->p_rts_flags |= (f); \
75      } while(0)
76
77  /* Clear flag and enqueue if the process was not runnable but is now. */
78  #define RTS_UNSET(rp, f) \
79      do { \
80          int rts; \
81          rts = (rp)->p_rts_flags; \
82          (rp)->p_rts_flags &= ~(f); \
83          if(rts && !(rp)->p_rts_flags) { enqueue(rp); } \
84      } while(0)
85
86  /* Set flag and dequeue if the process was runnable. */
87  #define RTS_LOCK_SET(rp, f) \
88      do { \
89          if(!(rp)->p_rts_flags) { lock_dequeue(rp); } \
90          (rp)->p_rts_flags |= (f); \
91      } while(0)
92
93  /* Clear flag and enqueue if the process was not runnable but is now. */
94  #define RTS_LOCK_UNSET(rp, f) \
95      do { \
96          int rts; \
97          rts = (rp)->p_rts_flags; \
98          (rp)->p_rts_flags &= ~(f); \
99          if(rts && !(rp)->p_rts_flags) { lock_enqueue(rp); } \
100      } while(0)
101
102  /* Set flags to this value. */
103  #define RTS_LOCK_SETFLAGS(rp, f) \
104      do { \
105          if(!(rp)->p_rts_flags && (f)) { lock_dequeue(rp); } \
106          (rp)->p_rts_flags = (f); \
107      } while(0)
108
109  /* Misc flags */
110  #define REPLY_PENDING    0x01    /* reply to IPC_REQUEST is pending */
111  #define MF_VM            0x08    /* process uses VM */
112
113  /* Scheduling priorities for p_priority. Values must start at zero (highest
114   * priority) and increment. Priorities of the processes in the boot image
115   * can be set in table.c. IDLE must have a queue for itself, to prevent low
116   * priority user processes to run round-robin with IDLE.
117   */
118  #define NR_SCHED_QUEUES  16      /* MUST equal minimum priority + 1 */
119  #define TASK_Q           0      /* highest, used for kernel tasks */
120  #define MAX_USER_Q       0      /* highest priority for user processes */
121  #define USER_Q           7      /* default (should correspond to nice 0) */
122  #define MIN_USER_Q       14     /* minimum priority for user processes */
123  #define IDLE_Q           15     /* lowest, only IDLE process goes here */
124
125  /* Magic process table addresses. */
126  #define BEG_PROC_ADDR (&proc[0])
127  #define BEG_USER_ADDR (&proc[NR_TASKS])
128  #define END_PROC_ADDR (&proc[NR_TASKS + NR_PROCS])
129

```

```
130 #define NIL_PROC ((struct proc *) 0)
131 #define NIL_SYS_PROC ((struct proc *) 1)
132 #define cproc_addr(n) (&(proc + NR_TASKS)[(n)])
133 #define proc_addr(n) (pproc_addr + NR_TASKS)[(n)]
134 #define proc_nr(p) ((p)->p_nr)
135
136 #define isokprocn(n) ((unsigned) ((n) + NR_TASKS) < NR_PROCS + NR_TASKS)
137 #define isemptyn(n) isemptyp(proc_addr(n))
138 #define isemptyp(p) ((p)->p_rts_flags == SLOT_FREE)
139 #define iskernelp(p) iskerneln((p)->p_nr)
140 #define iskerneln(n) ((n) < 0)
141 #define isuserp(p) isusern((p)->p_nr)
142 #define isusern(n) ((n) >= 0)
143
144 /* The process table and pointers to process table slots. The pointers allow
145  * faster access because now a process entry can be found by indexing the
146  * pproc_addr array, while accessing an element i requires a multiplication
147  * with sizeof(struct proc) to determine the address.
148  */
149 EXTERN struct proc proc[NR_TASKS + NR_PROCS]; /* process table */
150 EXTERN struct proc *pproc_addr[NR_TASKS + NR_PROCS];
151 EXTERN struct proc *rdy_head[NR_SCHED_QUEUES]; /* ptrs to ready list headers */
152 EXTERN struct proc *rdy_tail[NR_SCHED_QUEUES]; /* ptrs to ready list tails */
153
154 #endif /* PROC_H */
155
```