

```

1  /* The kernel call implemented in this file:
2  *   m_type:   SYS_FORK
3  *
4  * The parameters for this kernel call are:
5  *   ml_i1:    PR_SLOT (child's process table slot)
6  *   ml_i2:    PR_ENDPT (parent, process that forked)
7  */
8
9  #include "../system.h"
10 #include <signal.h>
11
12 #include <minix/endpoint.h>
13
14 #if USE_FORK
15
16 /*=====
17 *                                     do_fork                                     *
18 *=====*/
19 PUBLIC int do_fork(m_ptr)
20 register message *m_ptr;          /* pointer to request message */
21 {
22     /* Handle sys_fork(). PR_ENDPT has forked. The child is PR_SLOT. */
23     #if (_MINIX_CHIP == _CHIP_INTEL)
24         reg_t old_ldt_sel;
25     #endif
26     register struct proc *rp;      /* process pointer */
27     register struct proc *rpc;     /* child process pointer */
28     struct proc *rpp;              /* parent process pointer */
29     struct mem_map *map_ptr;       /* virtual address of map inside caller (PM) */
30     int i, gen, r;
31     int p_proc;
32
33     if(!isokendpt(m_ptr->PR_ENDPT, &p_proc))
34         return EINVAL;
35     rpp = proc_addr(p_proc);
36     rpc = proc_addr(m_ptr->PR_SLOT);
37     if (isemptyp(rpp) || ! isemptyp(rpc)) return(EINVAL);
38
39     map_ptr= (struct mem_map *) m_ptr->PR_MEM_PTR;
40
41     /* Copy parent 'proc' struct to child. And reinitialize some fields. */
42     gen = _ENDPOINT_G(rpc->p_endpoint);
43     #if (_MINIX_CHIP == _CHIP_INTEL)
44         old_ldt_sel = rpc->p_seg.p_ldt_sel; /* backup local descriptors */
45         *rpc = *rpp;                       /* copy 'proc' struct */
46         rpc->p_seg.p_ldt_sel = old_ldt_sel; /* restore descriptors */
47     #else
48         *rpc = *rpp;                       /* copy 'proc' struct */
49     #endif
50     if(++gen >= _ENDPOINT_MAX_GENERATION) /* increase generation */
51         gen = 1;                          /* generation number wraparound */
52     rpc->p_nr = m_ptr->PR_SLOT;             /* this was obliterated by copy */
53     rpc->p_endpoint = _ENDPOINT(gen, rpc->p_nr); /* new endpoint of slot */
54
55     rpc->p_reg.retreg = 0;                  /* child sees pid = 0 to know it is child */
56     rpc->p_user_time = 0;                   /* set all the accounting times to 0 */
57     rpc->p_sys_time = 0;
58
59     /* Because this is a copy of the parent process, message data is copied over
60     * as well. This should be reset so we have a clean slate.
61     */
62     memset(rpc->p_mess_sent, 0, sizeof(rpc->p_mess_sent));
63
64     /* Reset the number of messages sent by other processes to any previous
65     * process that used the same pid.

```

```
66      */
67      for (rp = BEG_PROC_ADDR, i = rpc->p_nr + NR_TASKS; rp < END_PROC_ADDR; ++rp)
68          rp->p_mess_sent[i] = 0;
69
70      /* Parent and child have to share the quantum that the forked process had,
71       * so that queued processes do not have to wait longer because of the fork.
72       * If the time left is odd, the child gets an extra tick.
73       */
74      rpc->p_ticks_left = (rpc->p_ticks_left + 1) / 2;
75      rpp->p_ticks_left = rpp->p_ticks_left / 2;
76
77      /* If the parent is a privileged process, take away the privileges from the
78       * child process and inhibit it from running by setting the NO_PRIV flag.
79       * The caller should explicitly set the new privileges before executing.
80       */
81      if (priv(rpp)->s_flags & SYS_PROC) {
82          rpc->p_priv = priv_addr(USER_PRIV_ID);
83          rpc->p_rts_flags |= NO_PRIV;
84      }
85
86      /* Calculate endpoint identifier, so caller knows what it is. */
87      m_ptr->PR_ENDPT = rpc->p_endpoint;
88
89      /* Install new map */
90      r = newmap(rpc, map_ptr);
91
92      /* Only one in group should have SIGNALED, child doesn't inherit tracing. */
93      RTS_LOCK_UNSET(rpc, (SIGNALED | SIG_PENDING | P_STOP));
94      sigemptyset(&rpc->p_pending);
95
96      return r;
97  }
98
99  #endif /* USE_FORK */
100
101
```