

```

586     check_runqueues("dequeue2");
587 #endif
588 }
589
590 /*=====
591 *                                sched                                *
592 *=====*/
593 PRIVATE void sched(rp, queue, front)
594 register struct proc *rp;          /* process to be scheduled */
595 int *queue;                        /* return: queue to use */
596 int *front;                        /* return: front or back */
597 {
598 /* This function determines the scheduling policy. It is called whenever a
599 * process must be added to one of the scheduling queues to decide where to
600 * insert it. As a side-effect the process' priority may be updated.
601 */
602     int time_left = (rp->p_ticks_left > 0);    /* quantum fully consumed */
603
604     /* Check whether the process has time left. Otherwise give a new quantum
605     * and lower the process' priority, unless the process already is in the
606     * lowest queue.
607     */
608     if (! time_left) {                  /* quantum consumed ? */
609         rp->p_ticks_left = rp->p_quantum_size; /* give new quantum */
610         if (rp->p_priority < (IDLE_Q-1)) {
611             rp->p_priority += 1;        /* lower priority */
612         }
613     }
614
615     /* If there is time left, the process is added to the front of its queue,
616     * so that it can immediately run. The queue to use simply is always the
617     * process' current priority.
618     */
619     *queue = rp->p_priority;
620     *front = time_left;
621 }
622
623 /*=====
624 *                                pick_proc                            *
625 *=====*/
626 PRIVATE void pick_proc()
627 {
628 /* Decide who to run now. A new process is selected by setting 'next_ptr'.
629 * When a billable process is selected, record it in 'bill_ptr', so that the
630 * clock task can tell who to bill for system time.
631 */
632     register struct proc *rp;          /* process to run */
633     int q;                             /* iterate over queues */
634     register struct proc *low_rp = NIL_PROC; /* lowest time user process */
635     clock_t low_time = LONG_MAX;       /* lowest time */
636
637     /* Check each of the scheduling queues except for the idle queue for ready
638     * processes. The number of queues is defined in proc.h, and priorities are
639     * set in the task table.
640     */
641     for (q = 0; q < NR_SCHED_QUEUES - 1; q++) {
642         for (rp = rdy_head[q]; rp != NIL_PROC; rp = rp->p_nextready) {
643             /* System process. */
644             if (priv(rp)->s_flags & SYS_PROC) {
645                 next_ptr = rp;        /* run process 'rp' next */
646                 return;
647             }
648             /* User process. */
649             if (rp->p_recent_time < low_time) {
650                 low_time = rp->p_recent_time; /* record lowest */

```

```

651         low_rp = rp;                                /* time and proc */
652     }
653 }
654 }
655 /* The function has scanned through the scheduling queues and determined
656  * that no system tasks wish to run while also recording the user process
657  * (if any) with the lowest recent CPU time. Now test to see if there ARE
658  * any user processes that wish to run.
659  */
660 if (low_rp != NIL_PROC) {
661     next_ptr = low_rp;                                /* run process 'low_rp' next */
662     bill_ptr = low_rp;                                /* bill for system time */
663     return;
664 }
665 /* No other process wants to run, so attempt to run the idle process. */
666 if ( (rp = rdy_head[q]) != NIL_PROC) {
667     next_ptr = rp;                                    /* run process 'rp' next */
668     bill_ptr = rp;                                    /* bill for system time */
669     return;
670 }
671 panic("no ready process", NO_NUM);
672 }
673
674 /*=====
675  *                                balance_queues                                *
676  *=====*/
677 #define Q_BALANCE_TICKS 100
678 PUBLIC void balance_queues(tp)
679 timer_t *tp;                                /* watchdog timer pointer */
680 {
681 /* Check entire process table and give all process a higher priority. This
682  * effectively means giving a new quantum. If a process already is at its
683  * maximum priority, its quantum will be renewed.
684  */
685     static timer_t queue_timer;                /* timer structure to use */
686     register struct proc* rp;                /* process table pointer */
687     clock_t next_period;                    /* time of next period */
688     int ticks_added = 0;                    /* total time added */
689
690     for (rp=BEG_PROC_ADDR; rp<END_PROC_ADDR; rp++) {
691         if (! isemptyp(rp)) {                /* check slot use */
692             rp->p_recent_time >= 1;            /* decay recent cpu */
693             lock(5, "balance_queues");
694             if (rp->p_priority > rp->p_max_priority) { /* update priority? */
695                 if (rp->p_rts_flags == 0) dequeue(rp); /* take off queue */
696                 ticks_added += rp->p_quantum_size; /* do accounting */
697                 rp->p_priority -= 1;            /* raise priority */
698                 if (rp->p_rts_flags == 0) enqueue(rp); /* put on queue */
699             }
700             else {
701                 ticks_added += rp->p_quantum_size - rp->p_ticks_left;
702                 rp->p_ticks_left = rp->p_quantum_size; /* give new quantum */
703             }
704             unlock(5);
705         }
706     }
707 #if DEBUG
708     kprintf("ticks_added: %d\n", ticks_added);
709 #endif
710
711 /* Now schedule a new watchdog timer to balance the queues again. The
712  * period depends on the total amount of quantum ticks added.
713  */
714     next_period = MAX(Q_BALANCE_TICKS, ticks_added); /* calculate next */
715     set_timer(&queue_timer, get_uptime() + next_period, balance_queues);

```