

```

1  #ifndef PROC_H
2  #define PROC_H
3
4  /* Here is the declaration of the process table.  It contains all process
5   * data, including registers, flags, scheduling priority, memory map,
6   * accounting, message passing (IPC) information, and so on.
7   *
8   * Many assembly code routines reference fields in it.  The offsets to these
9   * fields are defined in the assembler include file sconst.h.  When changing
10  * struct proc, be sure to change sconst.h to match.
11  */
12  #include <minix/com.h>
13  #include "const.h"
14  #include "priv.h"
15
16  struct proc {
17      struct stackframe_s p_reg;      /* process' registers saved in stack frame */
18      struct segframe p_seg;          /* segment descriptors */
19      proc_nr_t p_nr;                 /* number of this process (for fast access) */
20      struct priv *p_priv;             /* system privileges structure */
21      short p_rts_flags;              /* process is runnable only if zero */
22      short p_misc_flags;             /* flags that do not suspend the process */
23
24      char p_priority;                /* current scheduling priority */
25      char p_max_priority;            /* maximum scheduling priority */
26      char p_ticks_left;              /* number of scheduling ticks left */
27      char p_quantum_size;            /* quantum size in ticks */
28
29      struct mem_map p_memmap[NR_LOCAL_SEGS]; /* memory map (T, D, S) */
30
31      clock_t p_user_time;            /* user time in ticks */
32      clock_t p_sys_time;             /* sys time in ticks */
33      clock_t p_recent_time;          /* recent time in ticks */
34
35      struct proc *p_nextready;       /* pointer to next ready process */
36      struct proc *p_caller_q;        /* head of list of procs wishing to send */
37      struct proc *p_q_link;          /* link to next proc wishing to send */
38      message *p_messbuf;             /* pointer to passed message buffer */
39      int p_getfrom_e;                /* from whom does process want to receive? */
40      int p_sendto_e;                /* to whom does process want to send? */
41
42      sigset_t p_pending;             /* bit map for pending kernel signals */
43
44      char p_name[P_NAME_LEN];        /* name of the process, including \0 */
45
46      endpoint_t p_endpoint;          /* endpoint number, generation-aware */
47
48      #if DEBUG_SCHED_CHECK
49          int p_ready, p_found;
50      #endif
51  };
52
53  /* Bits for the runtime flags.  A process is runnable iff p_rts_flags == 0. */
54  #define SLOT_FREE      0x01  /* process slot is free */
55  #define NO_PRIORITY    0x02  /* process has been stopped */
56  #define SENDING        0x04  /* process blocked trying to send */
57  #define RECEIVING      0x08  /* process blocked trying to receive */
58  #define SIGNALLED      0x10  /* set when new kernel signal arrives */
59  #define SIG_PENDING    0x20  /* unready while signal being processed */
60  #define P_STOP         0x40  /* set when process is being traced */
61  #define NO_PRIV        0x80  /* keep forked system process from running */
62  #define NO_ENDPOINT    0x100 /* process cannot send or receive messages */
63
64  /* These runtime flags can be tested and manipulated by these macros. */
65

```

```

1  /* This file contains the main program of MINIX as well as its shutdown code.
2  * The routine main() initializes the system and starts the ball rolling by
3  * setting up the process table, interrupt vectors, and scheduling each task
4  * to run to initialize itself.
5  * The routine shutdown() does the opposite and brings down MINIX.
6  *
7  * The entries into this file are:
8  *   main:             MINIX main program
9  *   prepare_shutdown: prepare to take MINIX down
10 *
11 #include "kernel.h"
12 #include <signal.h>
13 #include <string.h>
14 #include <unistd.h>
15 #include <a.out.h>
16 #include <minix/callnr.h>
17 #include <minix/com.h>
18 #include <minix/endpoint.h>
19 #include "proc.h"
20
21 /* Prototype declarations for PRIVATE functions. */
22 FORWARD _PROTOTYPE( void announce, (void));
23 FORWARD _PROTOTYPE( void shutdown, (timer_t *));
24
25 /*=====*
26 *                               main                               *
27 *=====*/
28 PUBLIC void main()
29 {
30 /* Start the ball rolling. */
31 struct boot_image *ip;      /* boot image pointer */
32 register struct proc *rp;    /* process pointer */
33 register struct priv *sp;    /* privilege structure pointer */
34 register int i, s;
35 int hdrindex;               /* index to array of a.out headers */
36 phys_clicks text_base;
37 vir_clicks text_clicks, data_clicks, st_clicks;
38 reg_t ktsb;                 /* kernel task stack base */
39 struct exec e_hdr;          /* for a copy of an a.out header */
40
41 /* Clear the process table. Anounce each slot as empty and set up mappings
42 * for proc_addr() and proc_nr() macros. Do the same for the table with
43 * privilege structures for the system processes.
44 */
45 for (rp = BEG_PROC_ADDR, i = -NR_TASKS; rp < END_PROC_ADDR; ++rp, ++i) {
46     rp->p_rts_flags = SLOT_FREE;      /* initialize free slot */
47     rp->p_nr = i;                     /* proc number from ptr */
48     rp->p_endpoint = _ENDPOINT(0, rp->p_nr); /* generation no. 0 */
49     (pproc_addr + NR_TASKS)[i] = rp;  /* proc ptr from number */
50     rp->p_recent_time = 0;             /* recent cpu time */
51 }
52 for (sp = BEG_PRIV_ADDR, i = 0; sp < END_PRIV_ADDR; ++sp, ++i) {
53     sp->s_proc_nr = NONE;              /* initialize as free */
54     sp->s_id = i;                     /* priv structure index */
55     ppriv_addr[i] = sp;               /* priv ptr from number */
56 }
57
58 /* Set up proc table entries for processes in boot image. The stacks of the
59 * kernel tasks are initialized to an array in data space. The stacks
60 * of the servers have been added to the data segment by the monitor, so
61 * the stack pointer is set to the end of the data segment. All the
62 * processes are in low memory on the 8086. On the 386 only the kernel
63 * is in low memory, the rest is loaded in extended memory.
64 */
65

```

```

1  /* The kernel call implemented in this file:
2  *   m_type:      SYS_FORK
3  *
4  * The parameters for this kernel call are:
5  *   ml_i1:       PR_SLOT  (child's process table slot)
6  *   ml_i2:       PR_ENDPT (parent, process that forked)
7  */
8
9  #include "../system.h"
10 #include <signal.h>
11
12 #include <minix/endpoint.h>
13
14 #if USE_FORK
15
16 /*=====
17 *                                     do_fork                                     *
18 *=====*/
19 PUBLIC int do_fork(m_ptr)
20 register message *m_ptr;          /* pointer to request message */
21 {
22     /* Handle sys_fork().  PR_ENDPT has forked.  The child is PR_SLOT. */
23     #if (_MINIX_CHIP == _CHIP_INTEL)
24         reg_t old_ldt_sel;
25     #endif
26     register struct proc *rpc;          /* child process pointer */
27     struct proc *rpp;                   /* parent process pointer */
28     struct mem_map *map_ptr;            /* virtual address of map inside caller (PM) */
29     int i, gen, r;
30     int p_proc;
31
32     if(!isokendpt(m_ptr->PR_ENDPT, &p_proc))
33         return EINVAL;
34     rpp = proc_addr(p_proc);
35     rpc = proc_addr(m_ptr->PR_SLOT);
36     if (isempty(rpp) || ! isempty(rpc)) return(EINVAL);
37
38     map_ptr= (struct mem_map *) m_ptr->PR_MEM_PTR;
39
40     /* Copy parent 'proc' struct to child. And reinitialize some fields. */
41     gen = _ENDPOINT_G(rpc->p_endpoint);
42     #if (_MINIX_CHIP == _CHIP_INTEL)
43         old_ldt_sel = rpc->p_seg.p_ldt_sel;    /* backup local descriptors */
44         *rpc = *rpp;                          /* copy 'proc' struct */
45         rpc->p_seg.p_ldt_sel = old_ldt_sel;    /* restore descriptors */
46     #else
47         *rpc = *rpp;                          /* copy 'proc' struct */
48     #endif
49     if(++gen >= _ENDPOINT_MAX_GENERATION) /* increase generation */
50         gen = 1;                          /* generation number wraparound */
51     rpc->p_nr = m_ptr->PR_SLOT;              /* this was obliterated by copy */
52     rpc->p_endpoint = _ENDPOINT(gen, rpc->p_nr); /* new endpoint of slot */
53
54     rpc->p_reg.retreg = 0;                   /* child sees pid = 0 to know it is child */
55     rpc->p_user_time = 0;                   /* set all the accounting times to 0 */
56     rpc->p_sys_time = 0;
57     rpc->p_recent_time = 0;
58
59     /* Parent and child have to share the quantum that the forked process had,
60     * so that queued processes do not have to wait longer because of the fork.
61     * If the time left is odd, the child gets an extra tick.
62     */
63     rpc->p_ticks_left = (rpc->p_ticks_left + 1) / 2;
64     rpp->p_ticks_left = rpp->p_ticks_left / 2;
65

```

```

131 PRIVATE void init_clock()
132 {
133     /* First of all init the clock system.
134     *
135     * Here the (a) clock is set to produce a interrupt at
136     * every 1/60 second (ea. 60Hz).
137     *
138     * Running right away.
139     */
140     arch_init_clock(); /* architecture-dependent initialization. */
141
142     /* Initialize the CLOCK's interrupt hook. */
143     clock_hook.proc_nr_e = CLOCK;
144
145     put_irq_handler(&clock_hook, CLOCK_IRQ, clock_handler);
146     enable_irq(&clock_hook); /* ready for clock interrupts */
147
148     /* Set a watchdog timer to periodically balance the scheduling queues. */
149     balance_queues(NULL); /* side-effect sets new timer */
150 }
151
152 /*=====*
153 *                                clock_handler                                *
154 *=====*/
155 PRIVATE int clock_handler(hook)
156 irq_hook_t *hook;
157 {
158     /* This executes on each clock tick (i.e., every time the timer chip generates
159     * an interrupt). It does a little bit of work so the clock task does not have
160     * to be called on every tick. The clock task is called when:
161     *
162     *     (1) the scheduling quantum of the running process has expired, or
163     *     (2) a timer has expired and the watchdog function should be run.
164     *
165     * Many global global and static variables are accessed here. The safety of
166     * this must be justified. All scheduling and message passing code acquires a
167     * lock by temporarily disabling interrupts, so no conflicts with calls from
168     * the task level can occur. Furthermore, interrupts are not reentrant, the
169     * interrupt handler cannot be bothered by other interrupts.
170     *
171     * Variables that are updated in the clock's interrupt handler:
172     *     lost_ticks:
173     *         Clock ticks counted outside the clock task. This for example
174     *         is used when the boot monitor processes a real mode interrupt.
175     *     realtime:
176     *         The current uptime is incremented with all outstanding ticks.
177     *     proc_ptr, bill_ptr:
178     *         These are used for accounting. It does not matter if proc.c
179     *         is changing them, provided they are always valid pointers,
180     *         since at worst the previous process would be billed.
181     */
182     register unsigned ticks;
183
184     /* Get number of ticks and update realtime. */
185     ticks = lost_ticks + 1;
186     lost_ticks = 0;
187     realtime += ticks;
188
189     /* Update user and system accounting times. Charge the current process for
190     * user time. If the current process is not billable, that is, if a non-user
191     * process is running, charge the billable process for system time as well.
192     * Thus the unbillable process' user time is the billable user's system time.
193     */
194
195     proc_ptr->p_user_time += ticks;

```

```

196     proc_ptr->p_recent_time += ticks;
197     if (priv(proc_ptr)->s_flags & PREEMPTIBLE) {
198         proc_ptr->p_ticks_left -= ticks;
199     }
200     if (! (priv(proc_ptr)->s_flags & BILLABLE)) {
201         bill_ptr->p_sys_time += ticks;
202         bill_ptr->p_ticks_left -= ticks;
203     }
204
205     /* Update load average. */
206     load_update();
207
208     /* Check if do_clocktick() must be called. Done for alarms and scheduling.
209     * Some processes, such as the kernel tasks, cannot be preempted.
210     */
211     if ((next_timeout <= realtime) || (proc_ptr->p_ticks_left <= 0)) {
212         prev_ptr = proc_ptr;                /* store running process */
213         lock_notify(HARDWARE, CLOCK);        /* send notification */
214     }
215     return(1);                               /* reenale interrupts */
216 }
217
218 /*=====
219 *                               get_uptime                               *
220 *=====*/
221 PUBLIC clock_t get_uptime(void)
222 {
223     /* Get and return the current clock uptime in ticks. */
224     return(realtime);
225 }
226
227 /*=====
228 *                               set_timer                               *
229 *=====*/
230 PUBLIC void set_timer(tp, exp_time, watchdog)
231 struct timer *tp;                /* pointer to timer structure */
232 clock_t exp_time;                /* expiration realtime */
233 tmr_func_t watchdog;            /* watchdog to be called */
234 {
235     /* Insert the new timer in the active timers list. Always update the
236     * next timeout time by setting it to the front of the active list.
237     */
238     tmrs_settimer(&clock_timers, tp, exp_time, watchdog, NULL);
239     next_timeout = clock_timers->tmr_exp_time;
240 }
241
242 /*=====
243 *                               reset_timer                               *
244 *=====*/
245 PUBLIC void reset_timer(tp)
246 struct timer *tp;                /* pointer to timer structure */
247 {
248     /* The timer pointed to by 'tp' is no longer needed. Remove it from both the
249     * active and expired lists. Always update the next timeout time by setting
250     * it to the front of the active list.
251     */
252     tmrs_clrtimer(&clock_timers, tp, NULL);
253     next_timeout = (clock_timers == NULL) ?
254         TMR_NEVER : clock_timers->tmr_exp_time;
255 }
256
257 /*=====
258 *                               load_update                               *
259 *=====*/
260 PRIVATE void load_update(void)

```

```

586     check_runqueues("dequeue2");
587 #endif
588 }
589
590 /*=====*
591 *                                sched                                *
592 *=====*/
593 PRIVATE void sched(rp, queue, front)
594 register struct proc *rp;          /* process to be scheduled */
595 int *queue;                       /* return: queue to use */
596 int *front;                       /* return: front or back */
597 {
598     /* This function determines the scheduling policy. It is called whenever a
599     * process must be added to one of the scheduling queues to decide where to
600     * insert it. As a side-effect the process' priority may be updated.
601     */
602     int time_left = (rp->p_ticks_left > 0);    /* quantum fully consumed */
603
604     /* Check whether the process has time left. Otherwise give a new quantum
605     * and lower the process' priority, unless the process already is in the
606     * lowest queue.
607     */
608     if (! time_left) {                  /* quantum consumed ? */
609         rp->p_ticks_left = rp->p_quantum_size; /* give new quantum */
610         if (rp->p_priority < (IDLE_Q-1)) {
611             rp->p_priority += 1;         /* lower priority */
612         }
613     }
614
615     /* If there is time left, the process is added to the front of its queue,
616     * so that it can immediately run. The queue to use simply is always the
617     * process' current priority.
618     */
619     *queue = rp->p_priority;
620     *front = time_left;
621 }
622
623 /*=====*
624 *                                pick_proc                            *
625 *=====*/
626 PRIVATE void pick_proc()
627 {
628     /* Decide who to run now. A new process is selected by setting 'next_ptr'.
629     * When a billable process is selected, record it in 'bill_ptr', so that the
630     * clock task can tell who to bill for system time.
631     */
632     register struct proc *rp;          /* process to run */
633     int q;                            /* iterate over queues */
634     register struct proc *low_rp = NIL_PROC; /* lowest time user process */
635     clock_t low_time = LONG_MAX;      /* lowest time */
636
637     /* Check each of the scheduling queues except for the idle queue for ready
638     * processes. The number of queues is defined in proc.h, and priorities are
639     * set in the task table.
640     */
641     for (q = 0; q < NR_SCHED_QUEUES - 1; q++) {
642         for (rp = rdy_head[q]; rp != NIL_PROC; rp = rp->p_nextready) {
643             /* System process. */
644             if (priv(rp)->s_flags & SYS_PROC) {
645                 next_ptr = rp;         /* run process 'rp' next */
646                 return;
647             }
648             /* User process. */
649             if (rp->p_recent_time < low_time) {
650                 low_time = rp->p_recent_time; /* record lowest */

```

```

651         low_rp = rp;                                /* time and proc */
652     }
653 }
654 }
655 /* The function has scanned through the scheduling queues and determined
656  * that no system tasks wish to run while also recording the user process
657  * (if any) with the lowest recent CPU time. Now test to see if there ARE
658  * any user processes that wish to run.
659  */
660 if (low_rp != NIL_PROC) {
661     next_ptr = low_rp;                                /* run process 'low_rp' next */
662     bill_ptr = low_rp;                                /* bill for system time */
663     return;
664 }
665 /* No other process wants to run, so attempt to run the idle process. */
666 if ( (rp = rdy_head[q]) != NIL_PROC) {
667     next_ptr = rp;                                    /* run process 'rp' next */
668     bill_ptr = rp;                                    /* bill for system time */
669     return;
670 }
671 panic("no ready process", NO_NUM);
672 }
673
674 /*=====
675  *                                balance_queues                                *
676  *=====*/
677 #define Q_BALANCE_TICKS 100
678 PUBLIC void balance_queues(tp)
679 timer_t *tp;                                         /* watchdog timer pointer */
680 {
681 /* Check entire process table and give all process a higher priority. This
682  * effectively means giving a new quantum. If a process already is at its
683  * maximum priority, its quantum will be renewed.
684  */
685     static timer_t queue_timer;                      /* timer structure to use */
686     register struct proc* rp;                        /* process table pointer */
687     clock_t next_period;                             /* time of next period */
688     int ticks_added = 0;                             /* total time added */
689
690     for (rp=BEG_PROC_ADDR; rp<END_PROC_ADDR; rp++) {
691         if (! isemptyp(rp)) {                        /* check slot use */
692             rp->p_recent_time >= 1;                  /* decay recent cpu */
693             lock(5, "balance_queues");
694             if (rp->p_priority > rp->p_max_priority) { /* update priority? */
695                 if (rp->p_rts_flags == 0) dequeue(rp); /* take off queue */
696                 ticks_added += rp->p_quantum_size;    /* do accounting */
697                 rp->p_priority -= 1;                  /* raise priority */
698                 if (rp->p_rts_flags == 0) enqueue(rp); /* put on queue */
699             }
700             else {
701                 ticks_added += rp->p_quantum_size - rp->p_ticks_left;
702                 rp->p_ticks_left = rp->p_quantum_size; /* give new quantum */
703             }
704             unlock(5);
705         }
706     }
707 #if DEBUG
708     kprintf("ticks_added: %d\n", ticks_added);
709 #endif
710
711     /* Now schedule a new watchdog timer to balance the queues again. The
712     * period depends on the total amount of quantum ticks added.
713     */
714     next_period = MAX(Q_BALANCE_TICKS, ticks_added); /* calculate next */
715     set_timer(&queue_timer, get_uptime() + next_period, balance_queues);

```

```
1  /* Function prototypes. */
2
3  /* main.c */
4  _PROTOTYPE( int  main, (int argc, char **argv)           );
5
6  /* dmp.c */
7  _PROTOTYPE( int  do_fkey_pressed, (message *m)           );
8  _PROTOTYPE( void mapping_dmp, (void)                     );
9
10 /* dmp_kernel.c */
11 _PROTOTYPE( void proctab_dmp, (void)                      );
12 _PROTOTYPE( void memmap_dmp, (void)                      );
13 _PROTOTYPE( void privileges_dmp, (void)                  );
14 _PROTOTYPE( void sendmask_dmp, (void)                    );
15 _PROTOTYPE( void image_dmp, (void)                      );
16 _PROTOTYPE( void irqtab_dmp, (void)                      );
17 _PROTOTYPE( void kmessages_dmp, (void)                   );
18 _PROTOTYPE( void sched_dmp, (void)                      );
19 _PROTOTYPE( void monparams_dmp, (void)                   );
20 _PROTOTYPE( void kenv_dmp, (void)                       );
21 _PROTOTYPE( void timing_dmp, (void)                      );
22 _PROTOTYPE( void recent_dmp, (void)                      );
23
24 /* dmp_pm.c */
25 _PROTOTYPE( void mproc_dmp, (void)                      );
26 _PROTOTYPE( void sigaction_dmp, (void)                   );
27 _PROTOTYPE( void holes_dmp, (void)                      );
28
29 /* dmp_fs.c */
30 _PROTOTYPE( void dtab_dmp, (void)                       );
31 _PROTOTYPE( void fproc_dmp, (void)                      );
32
33 /* dmp_rs.c */
34 _PROTOTYPE( void rproc_dmp, (void)                      );
35
36 /* dmp_ds.c */
37 _PROTOTYPE( void data_store_dmp, (void)                  );
38
```



```

1  /* This file contains information dump procedures. During the initialization
2  * of the Information Service 'known' function keys are registered at the TTY
3  * server in order to receive a notification if one is pressed. Here, the
4  * corresponding dump procedure is called.
5  *
6  * The entry points into this file are
7  *   handle_fkey:      handle a function key pressed notification
8  */
9
10 #include "inc.h"
11
12 /* Define hooks for the debugging dumps. This table maps function keys
13  * onto a specific dump and provides a description for it.
14  */
15 #define NHOOKS 19
16
17 struct hook_entry {
18     int key;
19     void (*function)(void);
20     char *name;
21 } hooks[NHOOKS] = {
22     { F1,  proctab_dmp, "Kernel process table" },
23     { F2,  memmap_dmp, "Process memory maps" },
24     { F3,  image_dmp, "System image" },
25     { F4,  privileges_dmp, "Process privileges" },
26     { F5,  monparams_dmp, "Boot monitor parameters" },
27     { F6,  irqtab_dmp, "IRQ hooks and policies" },
28     { F7,  kmessages_dmp, "Kernel messages" },
29     { F9,  sched_dmp, "Scheduling queues" },
30     { F10, kenv_dmp, "Kernel parameters" },
31     { F11, timing_dmp, "Timing details (if enabled)" },
32     { F12, recent_dmp, "Recent CPU time" },
33     { SF1, mproc_dmp, "Process manager process table" },
34     { SF2, sigaction_dmp, "Signals" },
35     { SF3, fproc_dmp, "Filesystem process table" },
36     { SF4, dtab_dmp, "Device/Driver mapping" },
37     { SF5, mapping_dmp, "Print key mappings" },
38     { SF6, rproc_dmp, "Reincarnation server process table" },
39     { SF7, holes_dmp, "Memory free list" },
40     { SF8, data_store_dmp, "Data store contents" },
41 };
42
43 /*=====*
44  *                      handle_fkey                      *
45  *=====*/
46 #define pressed(k) ((F1<=(k)&&(k)<=F12 && bit_isset(m->FKEY_FKEYS, ((k)-F1+1)))\
47  || (SF1<=(k) && (k)<=SF12 && bit_isset(m->FKEY_SFKEYS, ((k)-SF1+1))))
48 PUBLIC int do_fkey_pressed(m)
49 message *m;                                /* notification message */
50 {
51     int s, h;
52
53     /* The notification message does not convey any information, other
54      * than that some function keys have been pressed. Ask TTY for details.
55      */
56     m->m_type = FKEY_CONTROL;
57     m->FKEY_REQUEST = FKEY_EVENTS;
58     if (OK != (s=sendrec(TTY_PROC_NR, m)))
59         report("IS", "warning, sendrec to TTY failed", s);
60
61     /* Now check which keys were pressed: F1-F12, SF1-SF12. */
62     for(h=0; h < NHOOKS; h++)
63         if(pressed(hooks[h].key))
64             hooks[h].function();
65

```

```

196         if (i % 8 == 7) ipc_to[++j] = ' ';
197     }
198     ipc_to[j] = '\0';
199     printf("%8s %4d  %s  %s  %3d %7lu %7lu  %s\n",
200         ip->proc_name, ip->proc_nr,
201         s_flags_str(ip->flags), s_traps_str(ip->trap_mask),
202         ip->priority, (long)ip->initial_pc, ip->stksize, ipc_to);
203 }
204 printf("\n");
205 }
206
207 /*=====
208 *                                     recent_dmp                                     *
209 *=====*/
210 PUBLIC void recent_dmp()
211 {
212     register struct proc *rp;
213     static struct proc *oldrp = BEG_PROC_ADDR;
214     register struct priv *sp;
215     int r, i, n = 0;
216
217     /* First obtain a fresh copy of the current process and system table. */
218     if ((r = sys_getprivtab(priv)) != OK) {
219         report("IS","warning: couldn't get copy of system privileges table", r);
220         return;
221     }
222     if ((r = sys_getproctab(proc)) != OK) {
223         report("IS","warning: couldn't get copy of process table", r);
224         return;
225     }
226
227     printf("\n-nr---id--name----flags-pri-privaddr-recent\n");
228
229     for (rp = oldrp; rp < END_PROC_ADDR; rp++) {
230         if (isempty(rp)) continue;
231         if (++n > 23) break;
232         if (proc_nr(rp) == IDLE)         printf("(%2d) ", proc_nr(rp));
233         else if (proc_nr(rp) < 0)        printf("[%2d] ", proc_nr(rp));
234         else                             printf(" %2d ", proc_nr(rp));
235         r = -1;
236         for (sp = &priv[0]; sp < &priv[NR_SYS_PROCS]; sp++)
237             if (sp->s_proc_nr == rp->p_nr) { r ++; break; }
238         if (r == -1 && ! (rp->p_rts_flags & SLOT_FREE)) {
239             sp = &priv[USER_PRIV_ID];
240         }
241         printf("(%02u) %-7.7s %s  %2d %4x      %6d",
242             sp->s_id, rp->p_name,
243             s_flags_str(sp->s_flags),
244             rp->p_priority,
245             rp->p_priv,
246             rp->p_recent_time
247         );
248         printf("\n");
249     }
250     if (rp == END_PROC_ADDR) rp = BEG_PROC_ADDR; else printf("--more--\r");
251     oldrp = rp;
252 }
253
254 /*=====
255 *                                     sched_dmp                                     *
256 *=====*/
257 PUBLIC void sched_dmp()
258 {
259     struct proc *rdy_head[NR_SCHED_QUEUES];
260     struct kinfo kinfo;

```